

```

import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

from google.colab import drive
drive.mount('/content/drive')

Mounted at /content/drive

# loading the dataset to a Pandas DataFrame
credit_card_data = pd.read_csv('/content/creditcard.csv')

# first 5 rows of the dataset
credit_card_data.head()

```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21
0	0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...	-0.018307
1	0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.225775
2	1	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.247998
3	1	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.108300
4	2	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	-0.009431

5 rows × 31 columns

```
credit_card_data.tail()
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21
25807	-0.769852	2.704375	-2.083145	1.018899	1.083598	-1.255315	1.242032	-0.525902	1.466585	...	-0.448296	
25808	-0.897475	0.963371	0.997351	0.329928	0.998766	-1.287190	0.713085	0.019353	-0.859152	...	0.118559	
25809	-0.377066	0.984515	0.988848	-0.261443	0.563332	0.197124	0.489867	0.281753	-0.543286	...	-0.234987	
25810	-0.353184	0.311241	1.586426	-1.515835	-0.636334	-0.836015	0.441214	-0.188933	1.218595	...	0.049963	
25810	0.827638	-0.539202	1.108173	1.532278	-0.950308	0.344304	-0.467828	0.217786	0.858742	...	NaN	

× 31 columns

```

# dataset informations
credit_card_data.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14595 entries, 0 to 14594
Data columns (total 31 columns):
#   Column  Non-Null Count  Dtype
---  -
0   Time    14595 non-null    int64
1   V1      14595 non-null    float64
2   V2      14595 non-null    float64
3   V3      14595 non-null    float64
4   V4      14595 non-null    float64

```

```

5   V5      14595 non-null float64
6   V6      14595 non-null float64
7   V7      14595 non-null float64
8   V8      14595 non-null float64
9   V9      14595 non-null float64
10  V10     14595 non-null float64
11  V11     14595 non-null float64
12  V12     14595 non-null float64
13  V13     14595 non-null float64
14  V14     14595 non-null float64
15  V15     14595 non-null float64
16  V16     14595 non-null float64
17  V17     14595 non-null float64
18  V18     14595 non-null float64
19  V19     14595 non-null float64
20  V20     14595 non-null float64
21  V21     14594 non-null float64
22  V22     14594 non-null float64
23  V23     14594 non-null float64
24  V24     14594 non-null float64
25  V25     14594 non-null float64
26  V26     14594 non-null float64
27  V27     14594 non-null float64
28  V28     14594 non-null float64
29  Amount  14594 non-null float64
30  Class   14594 non-null float64

```

```
dtypes: float64(30), int64(1)
```

```
memory usage: 3.5 MB
```

```
# checking the number of missing values in each column
credit_card_data.isnull().sum()
```

```

Time      0
V1        0
V2        0
V3        0
V4        0
V5        0
V6        0
V7        0
V8        0
V9        0
V10       0
V11       0
V12       0
V13       0
V14       0
V15       0
V16       0
V17       0
V18       0
V19       0
V20       0
V21       1
V22       1
V23       1
V24       1
V25       1
V26       1
V27       1
V28       1
Amount    1
Class     1
dtype: int64

```

```
# distribution of legit transactions & fraudulent transactions
credit_card_data['Class'].value_counts()
```

```
0.0    14533
1.0      61
Name: Class, dtype: int64
```

This Dataset is highly unblanced

0 --> Normal Transaction

1 --> fraudulent transaction

```
# separating the data for analysis
legit = credit_card_data[credit_card_data.Class == 0]
fraud = credit_card_data[credit_card_data.Class == 1]
```

```
print(legit.shape)
print(fraud.shape)
```

```
(14533, 31)
(61, 31)
```

```
# statistical measures of the data
legit.Amount.describe()
```

```
count    14533.000000
mean       64.065668
std       176.589083
min         0.000000
25%        5.550000
50%       15.950000
75%       52.990000
max      7712.430000
Name: Amount, dtype: float64
```

```
fraud.Amount.describe()
```

```
count      61.000000
mean       88.402295
std       297.522823
min         0.000000
25%         1.000000
50%         1.000000
75%         3.790000
max      1809.680000
Name: Amount, dtype: float64
```

```
# compare the values for both transactions
credit_card_data.groupby('Class').mean()
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	..
<b>Class</b>											
<b>0.0</b>	10777.717402	-0.215448	0.262816	0.879788	0.274916	-0.106953	0.136946	-0.128751	-0.019781	0.974505	
<b>1.0</b>	13325.934426	-5.118440	4.839489	-9.563219	6.621692	-3.266514	-2.089348	-6.722938	1.489551	-2.929718	

2 rows × 30 columns

Under-Sampling

Build a sample dataset containing similar distribution of normal transactions and Fraudulent Transactions

Number of Fraudulent Transactions --> 492

```
legit_sample = legit.sample(n=492)
```

Concatenating two DataFrames

```
new_dataset = pd.concat([legit_sample, fraud], axis=0)
```

```
new_dataset.head()
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	
<b>5111</b>	4784	1.374105	-0.253948	0.113605	0.069158	-0.052759	0.404988	-0.408880	-0.164304	0.450830	...	-0.6
<b>7770</b>	10823	0.625403	-1.140300	0.322045	0.473485	-0.836196	0.153164	-0.164060	0.058232	1.930849	...	-0.0
<b>6116</b>	7001	0.771466	-0.598544	2.246886	3.139235	-1.127466	2.269853	-1.620516	0.885450	2.449515	...	0.0
<b>13185</b>	23170	-3.743846	-3.265379	0.995384	3.041958	0.965919	-0.587442	1.862605	-0.117721	-0.399114	...	0.0
<b>10090</b>	15324	-3.690755	3.414470	0.714492	-1.985861	-0.543478	-1.366884	1.123768	-0.849273	5.116339	...	-0.9

5 rows × 31 columns

```
new_dataset.tail()
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	
<b>14104</b>	25095	1.192396	1.338974	-0.678876	3.123672	0.643245	-1.184323	0.397586	-0.253499	0.411135	...	
<b>14170</b>	25198	-15.903635	10.393917	-19.133602	6.185969	-12.538021	-4.027030	-13.897827	10.662252	-2.844954	...	
<b>14197</b>	25231	-16.598665	10.541751	-19.818982	6.017295	-13.025901	-4.128779	-14.118865	11.161144	-4.099551	...	
<b>14211</b>	25254	-17.275191	10.819665	-20.363886	6.046612	-13.465033	-4.166647	-14.409448	11.580797	-4.073856	...	
<b>14338</b>	25426	1.125336	1.130146	-0.962975	2.675688	0.990075	-0.243318	0.316192	0.122960	-1.143343	...	

5 rows × 31 columns

```
new_dataset['Class'].value_counts()
```

```
0.0    492
1.0     61
Name: Class, dtype: int64
```

```
new_dataset.groupby('Class').mean()
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	..
<b>Class</b>											
<b>0.0</b>	10255.290650	-0.228924	0.404879	0.835122	0.248904	-0.041809	0.170144	-0.116762	-0.053026	0.923775	
<b>1.0</b>	13325.934426	-5.118440	4.839489	-9.563219	6.621692	-3.266514	-2.089348	-6.722938	1.489551	-2.929718	

2 rows × 30 columns

## Splitting the data into Features &amp; Targets

```
X = new_dataset.drop(columns='Class', axis=1)
Y = new_dataset['Class']
```

```
print(X)
```

	Time	V1	V2	V3	V4	V5	V6	\
5111	4784	1.374105	-0.253948	0.113605	0.069158	-0.052759	0.404988	
7770	10823	0.625403	-1.140300	0.322045	0.473485	-0.836196	0.153164	
6116	7001	0.771466	-0.598544	2.246886	3.139235	-1.127466	2.269853	
13185	23170	-3.743846	-3.265379	0.995384	3.041958	0.965919	-0.587442	
10090	15324	-3.690755	3.414470	0.714492	-1.985861	-0.543478	-1.366884	
...	...	...	...	...	...	...	...	
14104	25095	1.192396	1.338974	-0.678876	3.123672	0.643245	-1.184323	
14170	25198	-15.903635	10.393917	-19.133602	6.185969	-12.538021	-4.027030	
14197	25231	-16.598665	10.541751	-19.818982	6.017295	-13.025901	-4.128779	
14211	25254	-17.275191	10.819665	-20.363886	6.046612	-13.465033	-4.166647	
14338	25426	1.125336	1.130146	-0.962975	2.675688	0.990075	-0.243318	
	V7	V8	V9	...	V20	V21	V22	\
5111	-0.408880	-0.164304	0.450830	...	-0.455944	-0.688367	-0.964160	
7770	-0.164060	0.058232	1.930849	...	0.360946	-0.025609	-0.280835	
6116	-1.620516	0.885450	2.449515	...	-0.242193	0.043859	0.629459	
13185	1.862605	-0.117721	-0.399114	...	2.299492	0.055303	-1.963346	
10090	1.123768	-0.849273	5.116339	...	2.370351	-0.965167	-0.590039	
...	...	...	...	...	...	...	...	
14104	0.397586	-0.253499	0.411135	...	-0.185455	-0.377503	-0.889597	
14170	-13.897827	10.662252	-2.844954	...	1.501565	1.577548	-1.280137	
14197	-14.118865	11.161144	-4.099551	...	1.534920	1.725853	-1.151606	
14211	-14.409448	11.580797	-4.073856	...	1.544970	1.729804	-1.208096	
14338	0.316192	0.122960	-1.143343	...	-0.138814	-0.166737	-0.521934	
	V23	V24	V25	V26	V27	V28	Amount	
5111	-0.087617	-0.971928	0.565894	0.418575	-0.006715	-0.002836	14.00	
7770	-0.269747	0.043434	0.176311	1.058848	-0.141906	0.027320	291.88	
6116	-0.097964	-0.327014	0.179940	0.251557	0.070432	0.025485	78.95	
13185	2.270957	-0.387028	0.510184	-0.487578	-0.377596	0.173759	788.77	
10090	0.046301	0.623774	0.368253	0.566924	0.708040	-0.355509	2.31	
...	...	...	...	...	...	...	...	
14104	-0.074208	0.035446	0.550578	-0.027171	-0.024921	0.073605	3.12	
14170	-0.601295	0.040404	0.995502	-0.273743	1.688136	0.527831	99.99	
14197	-0.680052	0.108176	1.066878	-0.233720	1.707521	0.511423	99.99	
14211	-0.726839	0.112540	1.119193	-0.233189	1.684063	0.503740	99.99	
14338	-0.112376	-0.592077	0.520791	0.043354	0.015159	0.063612	3.76	

```
[553 rows x 30 columns]
```

```
print(Y)
```

```
5111    0.0
7770    0.0
6116    0.0
13185   0.0
10090   0.0
...
14104   1.0
14170   1.0
14197   1.0
14211   1.0
14338   1.0
Name: Class, Length: 553, dtype: float64
```

## Split the data into Training data &amp; Testing Data

```
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, stratify=Y, random_state=2)

print(X.shape, X_train.shape, X_test.shape)

(553, 30) (442, 30) (111, 30)
```

## Model Training

### Logistic Regression

```
model = LogisticRegression()

# training the Logistic Regression Model with Training Data
model.fit(X_train, Y_train)

l/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge
AL NO. of ITERATIONS REACHED LIMIT.

the number of iterations (max_iter) or scale the data as shown in:
://scikit-learn.org/stable/modules/preprocessing.html
so refer to the documentation for alternative solver options:
://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression
i = _check_optimize_result(
.cRegression
Regression()
```

## Model Evaluation

### Accuracy Score

```
# accuracy on training data
X_train_prediction = model.predict(X_train)
training_data_accuracy = accuracy_score(X_train_prediction, Y_train)

print('Accuracy on Training data : ', training_data_accuracy)

Accuracy on Training data : 0.9909502262443439

# accuracy on test data
X_test_prediction = model.predict(X_test)
test_data_accuracy = accuracy_score(X_test_prediction, Y_test)

print('Accuracy score on Test Data : ', test_data_accuracy)

Accuracy score on Test Data : 0.972972972972973
```

